

```
# IMPORT DATA LIBRARIES import pandas as pd
```

Python - Linear Regression Model Cheat Sheet

```
import numpy as np # IMPORT VIS LIBRARIES
```

```
coeff_df = pd.DataFrame
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
%matplotlib inline
```

```
# IMPORT MODELLING LIBRARIES
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn import metrics
```

```
df = pd.read_csv('data.csv') read data df.head() check head df df.info()
```

```
check info df df.describe() check stats df df.columns check col names
```

```
sns.pairplot(df) pairplot sns.distplot(df['Y']) distribution plot
```

```
sns.heatmap(df.corr(), annot=True) heatmap with values
```

CREATE X and y -----

```
X = df[['col1','col2',etc.]] create df features y = df['col'] create df var to
```

```
predict ➡ SPLIT DATASET -----
```

pd.DataFrame: pd.DataFrame(data=None, index=None, columns=None, dtype=None, copy=False). **data** = values, **index**= name index, **columns**= name column. This could be useful just to interpret the

```
y,  
test_size=0.3)
```

```
X_train, X_test, y_train, y_test
```

```
= train_test_split(
```

```
X,
```

```
FIT THE MODEL -----  
split df in train and test df
```

```
lm.intercept_ show intercept
```

```
lm.coef_ show coefficients
```

```
(lm.coef_,X.columns,columns=['Coeff'])*
```

```
create coeff df
```

```
coefficient of the regression.
```

```
predictions = lm.predict(X_test) create predictions
```

```
plt.scatter(y_test,predictions)* plot predictions
```

```
sns.distplot((y_test-predictions),bins=50)* distplot of residuals
```

scatter: this graph show the difference between actual values and the values predicted by the model we trained. It should resemble as much as possible a **diagonal line**.

distplot: this graph shows the distributions of the residual errors, that is, the difference between the actual values minus the predicted values; it should result in an as much as possible **normal distribution**. If not, maybe change model!

```
print('MAE:', metrics.mean_absolute_error(y_test, predictions))
```

```
print('MSE:', metrics.mean_squared_error(y_test, predictions))
```

```
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test,  
predictions)))
```

MAE is the easiest to understand, because it's the average error. **MSE** is more popular than MAE, because MSE "punishes" larger errors, which tends to be useful in the real world.

RMSE is even more popular than MSE, because RMSE is interpretable in the "y" units.

```
lm = LinearRegression() instatiate model
```

```
lm.fit(X_train, y_train) train/fit the model ● SHOW
```

```
RESULTS -----
```